


# CS 455

# Database Management Systems



Department of Mathematics  
and Computer Science

Lecture 3  
Structured Query Language  
(SQL)

# Motivation

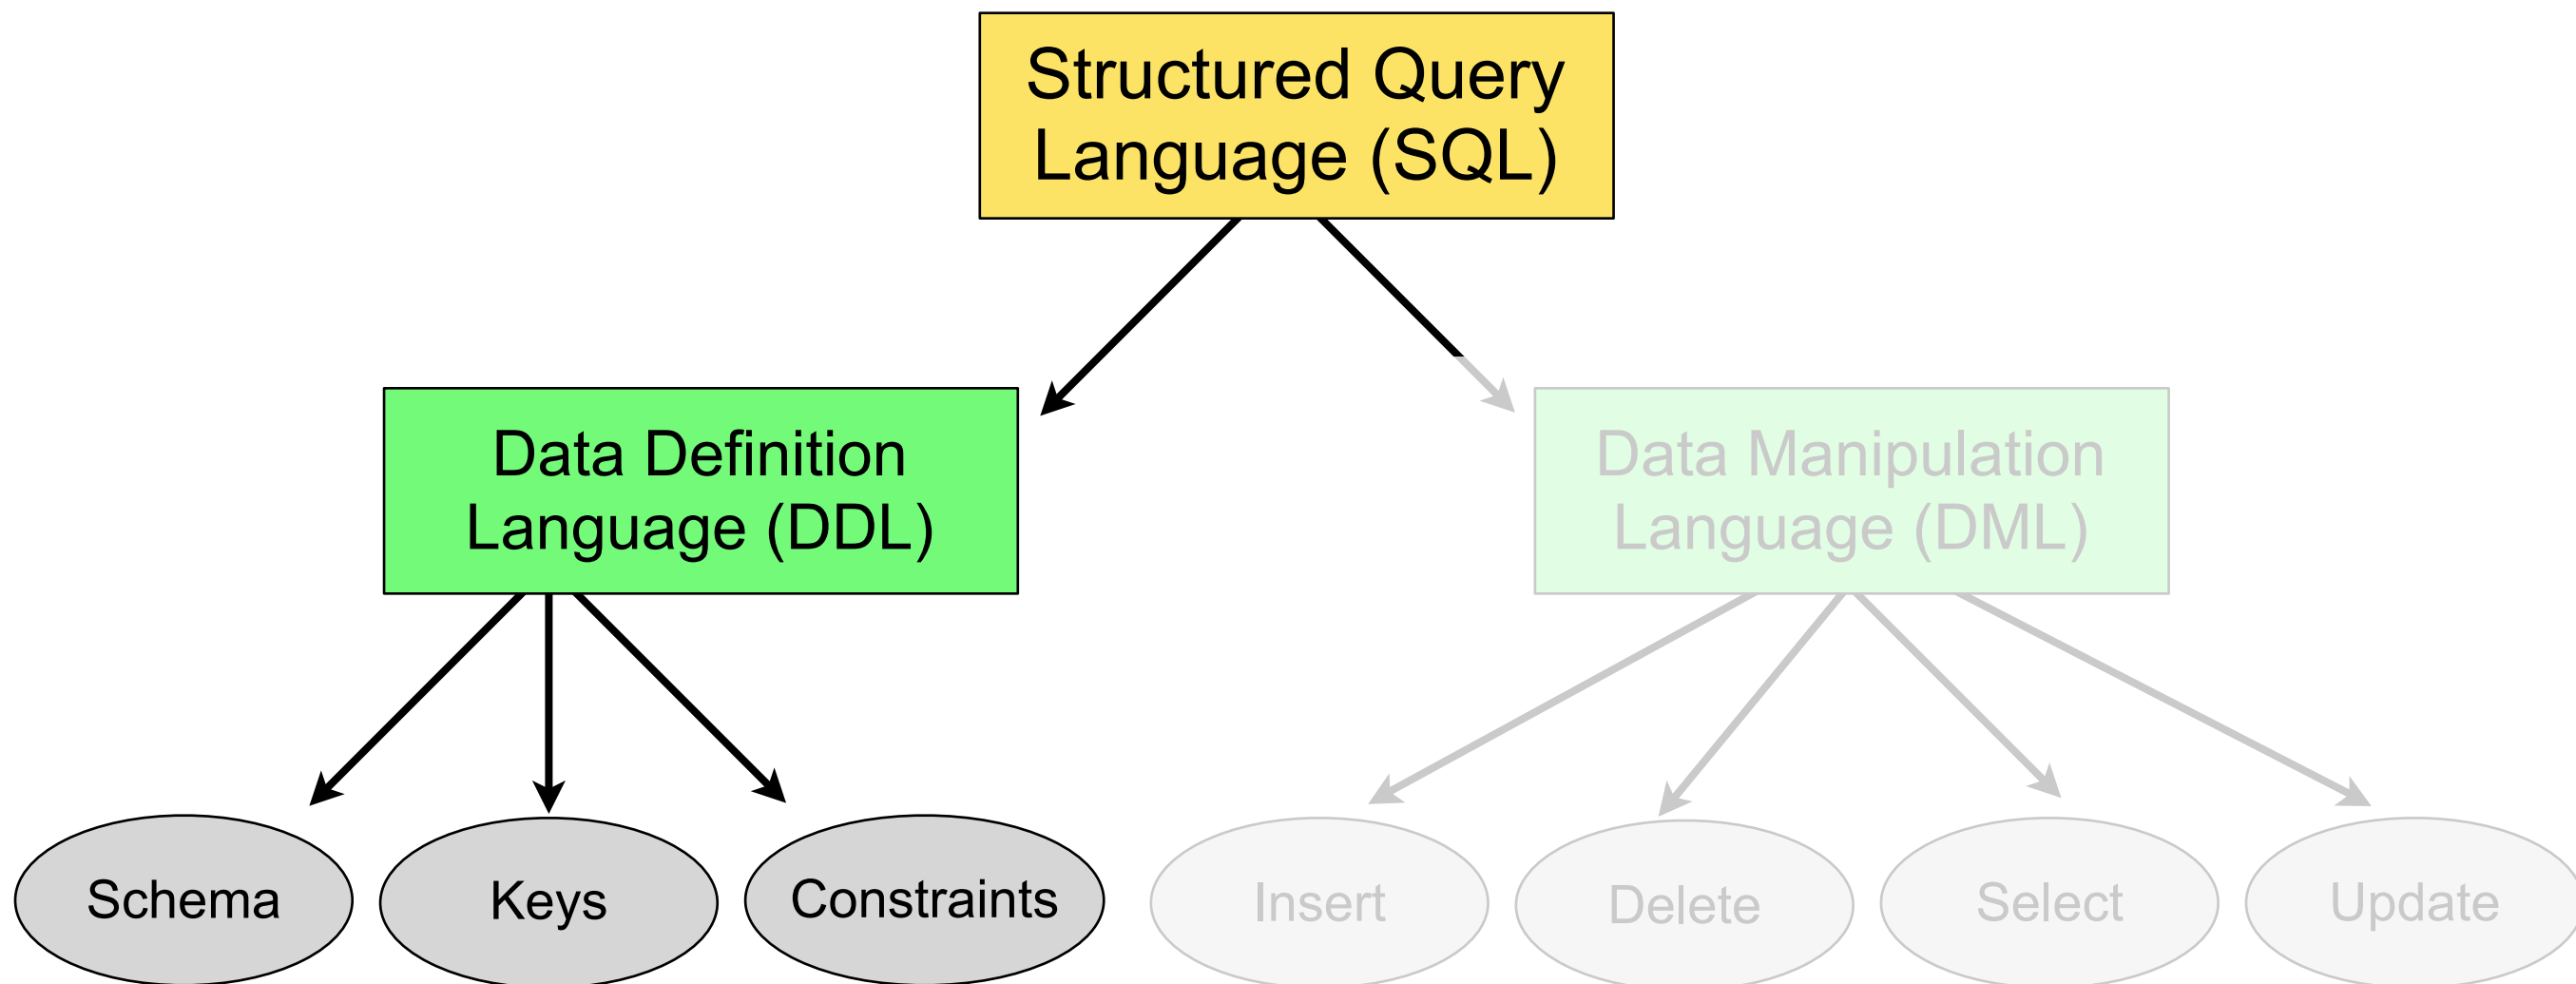
- ▶ Relational algebra is great, but...
  - DBMS are for use by common users, not just computer scientists
  - It's all "Greek" to lay users
$$\sigma, \Pi, \rho, \cup, \setminus, \times, \bowtie, \dots$$
  - Need user-friendly language, that has the same *expressivity* as relational algebra



# Topics

- ▶ Structured Query Language (SQL)
  - Data Definition Language (DDL)
    - Create table
    - Drop table
    - Alter table
  - Data Manipulation Language (DML)

# SQL at a Glance



*(In this class, we'll use SQLite3 syntax)  
... unfortunately, SQL flavors differ across implementations*

# Topics

- ▶ Structured Query Language (SQL)
  - Data Definition Language (DDL)
    - Create table
    - Drop table
    - Alter table
  - Data Manipulation Language (DML)

# Declaring a Relational Table

- ▶ Defining a Relation:  $R(a_1, a_2, \dots, a_n)$
- ▶ SQL Syntax:

```
create table R (  
    a_1 type_1    [c_1 c_2 ...][,    -- attribute 1  
    a_2 type_2    [c_1 c_2 ...],    -- attribute 2  
    .  
    .  
    a_n type_n    [c_1 c_2 ...],]    -- attribute n  
  
    [TC_1        -- table constraint 1  
    TC_2        -- table constraint 2  
    .  
    .  
    TC_k]       -- table constraint k  
);
```

*[...] means optional*

*-- is a line-comment*

# Attribute Data Types (in SQLite)

## ► INTEGER (Not INT!)

- Value is a signed integer
- 1 to 8 bytes (automatic: depending on magnitude of the number that's stored)

## ► REAL

- Value is a double-precision floating point number
- 8 bytes

## ► TEXT

- Value is a text string
- Stored as UTF8 (1-byte per char), UTF16 (2-bytes per char)

## ► BLOB (**B**inary **L**arge **O**bject)

- It could be an image, PDF, video, MP3, etc.

a\_1 type\_1 [c\_1 c\_2 ...]

Attribute name is followed by its data type

# Attribute Constraints

- ▶ Common Attribute (Column) Constraints: You can stack these constraints!

- NOT NULL
- UNIQUE
- CHECK(expression)
- PRIMARY KEY

Attribute Definition:

a\_1 type\_1 [c\_1 c\_2 ...]

Attribute constraints are optional, and stackable!

- If attribute type is an **integer**, will **auto-increment** if given **NULL** value.
- *(What if your primary key is a set of two or more attributes? See "Table Constraints")*

```
create table player (  
    pid      INTEGER PRIMARY KEY,  
    name     TEXT      UNIQUE,  
    salary   INTEGER NOT NULL CHECK(salary < 100000)  
);
```



# Example of SQLite Enforcing Constraints

```
create table player (  
  pid      INTEGER PRIMARY KEY,  
  name     TEXT      UNIQUE,  
  salary   INTEGER NOT NULL CHECK (salary < 100000)  
);
```

```
sqlite> INSERT INTO player VALUES (NULL, 'David', 25000);
```

```
sqlite> INSERT INTO player VALUES (1, 'Andy', 10000);
```

```
Error: PRIMARY KEY must be unique
```

```
sqlite> INSERT INTO player VALUES (NULL, 'David', 55000);
```

```
Error: column name is not unique
```

```
sqlite> INSERT INTO player VALUES (NULL, 'Tom', 55000);
```

```
sqlite> INSERT INTO player VALUES (NULL, 'Fred', 55000);
```

```
sqlite> INSERT INTO player VALUES (NULL, 'Jim', 150000);
```

```
Error: constraint failed
```

```
sqlite> select * from player;
```

```
1|David|25000
```

```
2|Tom|55000
```

```
3|Fred|55000
```


# Table Constraints

## ► Common Table Constraints

- PRIMARY KEY (*a\_1, ..., a\_n*)
- UNIQUE (*a\_1, ..., a\_n*)
- CHECK (*expression*)
- FOREIGN KEY (*a\_1, ..., a\_n*) REFERENCES *R'*(*b\_1, ..., b\_n*) ...

```
TC_1    -- table constraint 1
TC_2    -- table constraint 2
.
.
TC_k    -- table constraint k
```

Again, all optional and stackable!



*CASCADE* is super useful!!


*In SQLite, you have to enable foreign keys first!*

```
sqlite> PRAGMA foreign_keys = ON;
```

*Requires at least  
SQLite 3.6.19*

# Example Declaring Relational Tables

- Assume we have the following schema



Attribute constraints

```

create table player (
  pid    INTEGER PRIMARY KEY,
  name   TEXT     NOT NULL,
  salary INTEGER CHECK (salary < 100000)
);
    
```


Table constraints

```

create table playsFor (
  pid    INTEGER,
  team   TEXT,
  PRIMARY KEY (pid,team),
  FOREIGN KEY (pid) REFERENCES player(pid)
    ON UPDATE CASCADE
    ON DELETE CASCADE
);
    
```

# Foreign Key (Cascading In Action)

## ► Example of a Cascading Update



```
sqlite> PRAGMA foreign_keys = on;
... -- code to create and insert some data into player and playsFor tables
sqlite> select * from playsFor;
1|Blazers
2|Blazers

sqlite> select * from player;
1|David|25000
2|Tom|55000
3|Fred|55000

sqlite> update player set pid=5 where pid=1;
```

# Importance of Foreign Key Constraints

- What if the foreign key constraint *wasn't issued* when we created the table?

```
sqlite> select * from player;
```

```
1|David|55000
```

```
2|Tom|65000
```

```
3|Fred|75000
```

```
sqlite> select * from playsFor;
```

```
1|Blazers
```

```
2|Blazers
```

```
sqlite> update player set pid=6 where pid=2;
```

```
sqlite> select * from player;
```

```
1|David|55000
```

```
3|Fred|75000
```

```
6|Tom|65000
```

```
sqlite> select * from playsFor;
```

```
1|Blazers
```

```
2|Blazers -- stale!! Tom changed numbers! (Many problems now...)
```

# Other Useful DDL Commands

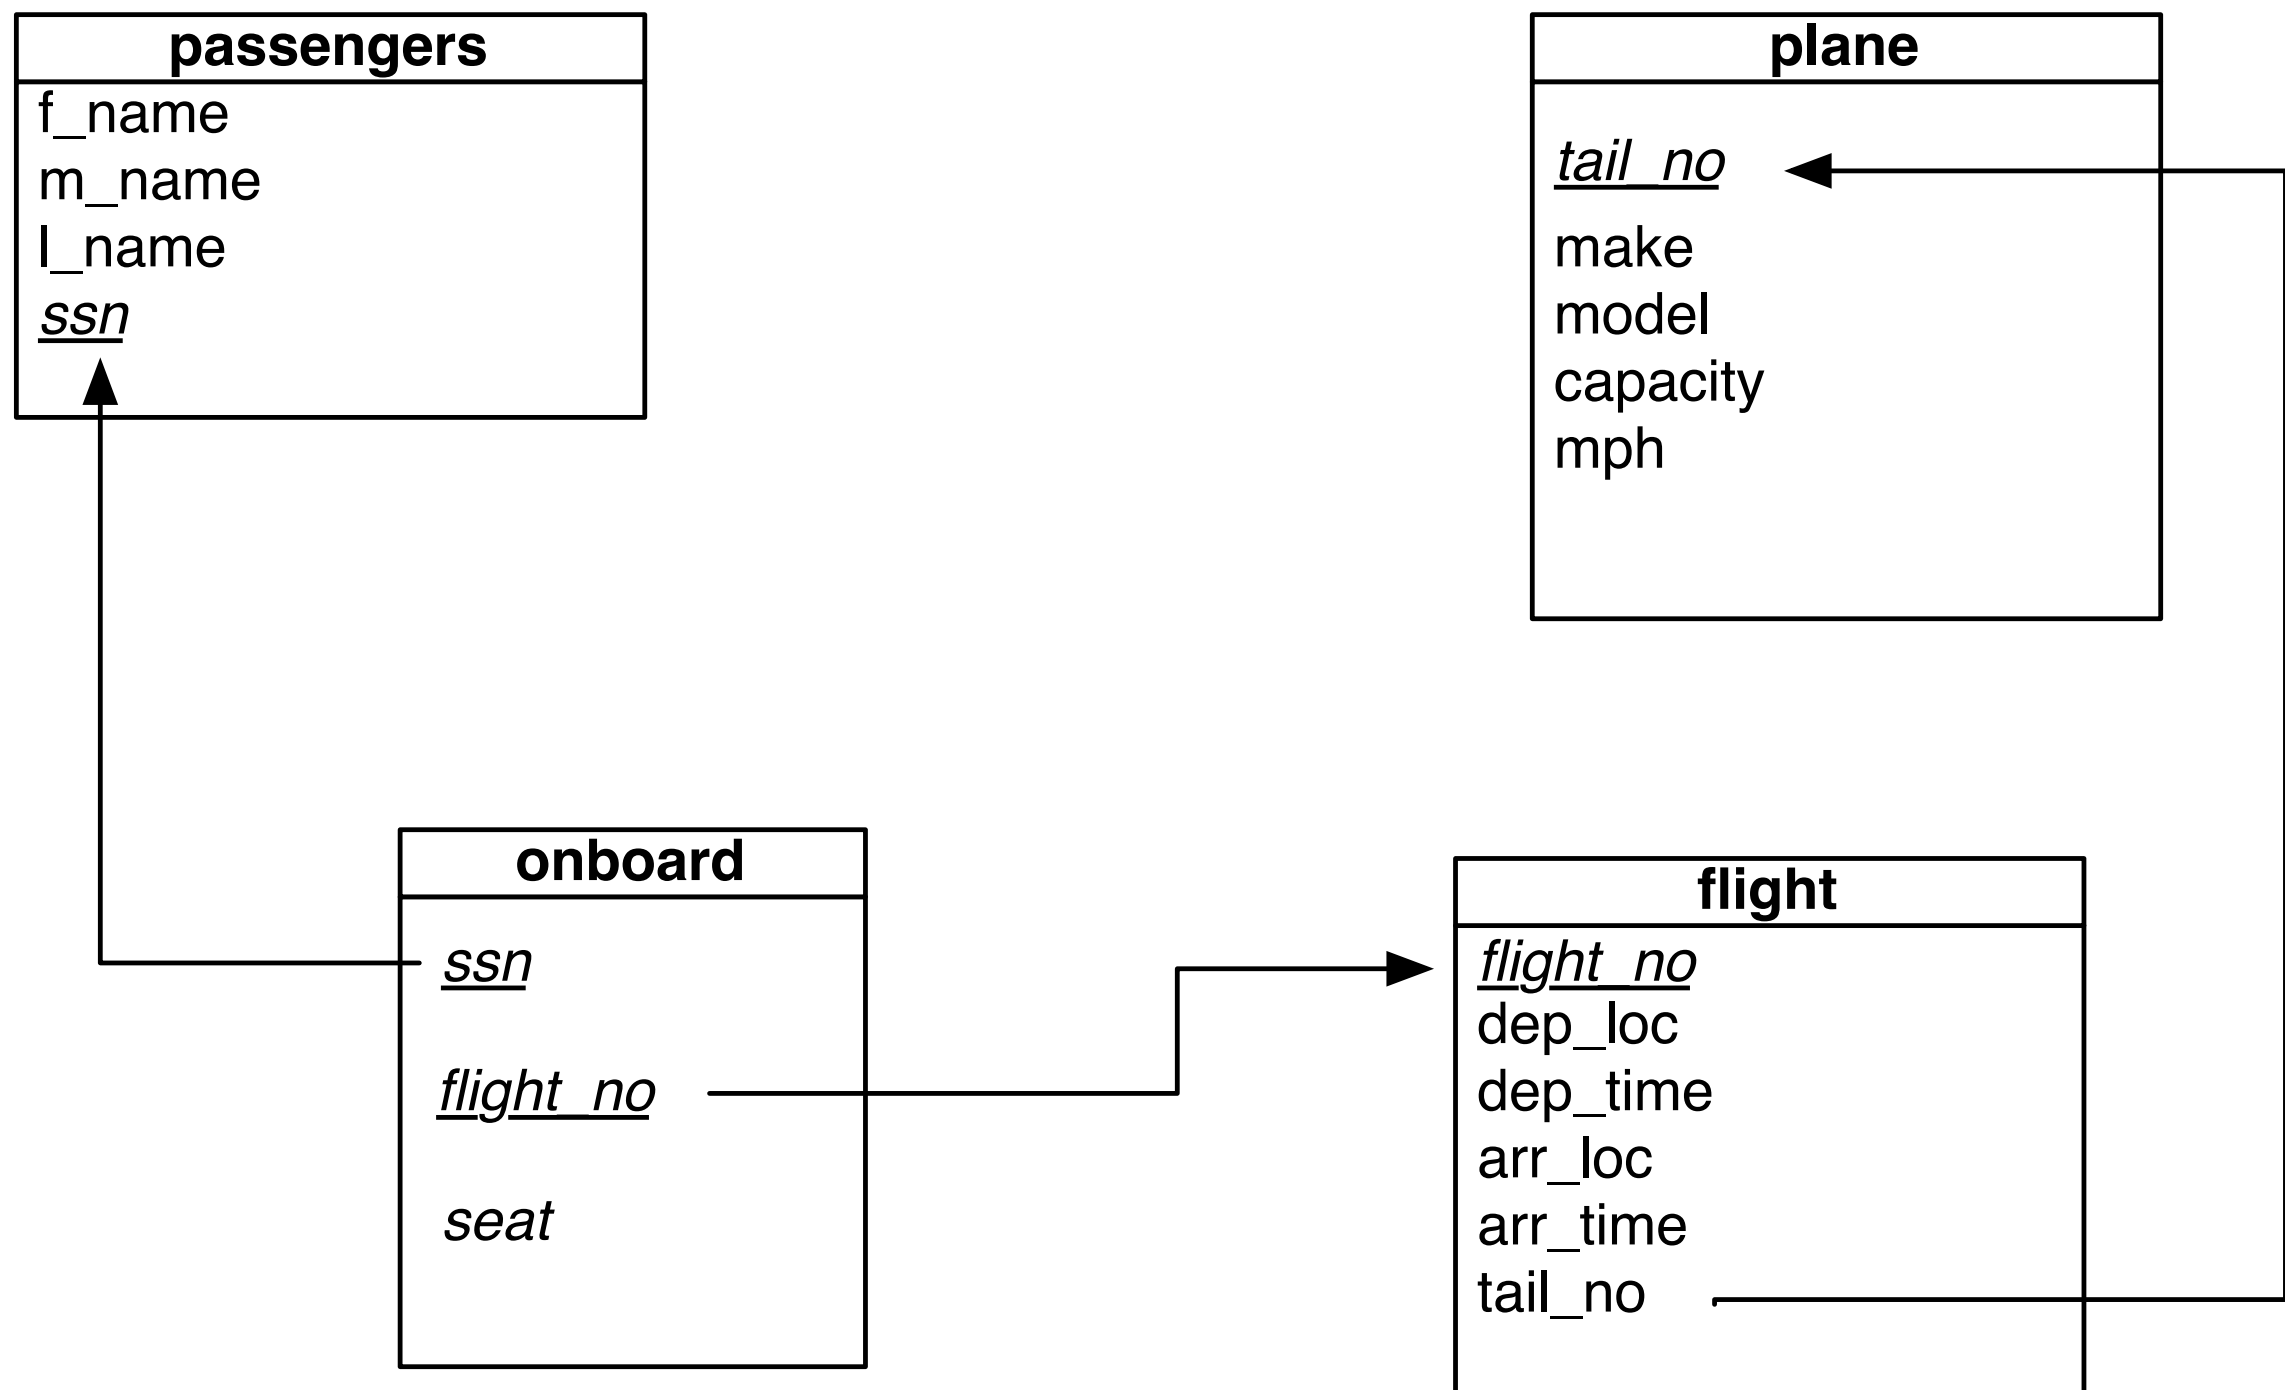
- ▶ Removing a relational table:

```
drop table [if exists] R;
```

- ▶ Changing the an existing table's properties: Look into [\*alter table\*](#) commands
  - For example, adding an attribute after-the-fact

```
alter table R  
  ADD attr datatype c1 c2, ...  
  FIRST|AFTER attr_name
```

# Create This Schema in SQLite!



(A solution on course page: [airport-schema.sql](#) after today's class)

# Topics

- ▶ Structured Query Language (SQL)
  - Data Definition Language (DDL)
  - Data Manipulation Language (DML)
    - Insert
    - Delete
    - Update
    - Select
      - from, where
      - order by
      - set operations
      - joins (implicit, explicit)
      - sub-queries
      - aggregation and grouping



# SQL: Inserting Tuples

► Relational Algebra Syntax:  $R \leftarrow R \cup E$

• Example:

$passengers \leftarrow passengers \cup \{('David', NULL, 'Chiu', '888-88-8888')\}$

► SQL Syntax: `INSERT INTO R VALUES (v1, v2, ..., vn);`

► Example:

```
INSERT INTO passengers VALUES
    ( 'David', NULL, 'Chiu', '888-88-8888' );
```

# SQL: Deleting Tuples

► Relational Algebra Syntax:  $R \leftarrow R \setminus E$

► Example:

$player \leftarrow player - \sigma_{Salary > 60000}(player)$

► SQL Syntax:

```
DELETE FROM R [WHERE C];
```

```
-- everything in [...] is optional  
-- That means WHERE clause is optional!  
-- C is assumed true if not given.
```

► Example:

```
DELETE FROM player WHERE Salary > 60000;
```

```
DELETE FROM player; -- OMG What did I just do??
```

# SQL: Updating Tuples

► Relational Algebra Syntax:  $R \leftarrow \Pi_{E_1, \dots, E_k}(R)$

► SQL Syntax:

```
UPDATE R SET a_1=e_1[, ..., a_k=e_k] [WHERE C];
```

*-- e\_i are SQL expressions*

*-- C is assumed true if not given*

► Example:

```
UPDATE passengers SET m_name='John' WHERE ssn='888-88-8888';
```

```
UPDATE players SET salary=(1.04*salary); -- I meant to leave off  
where clause..
```

# Select-From-Where (SFW)

- ▶ Recall this common R.A. expression:  $\Pi_{a_1, \dots, a_k}(\sigma_c(R))$
- ▶ In SQL, this projection-selection takes on a very common form:
  - Often referred to as an *SFW-query*

```
SELECT [distinct] a_1[,a_2,...,a_k] | *  
FROM R1[,R2,...,Rn]  
[WHERE C];
```

-- Once again, C is assumed true if not given

-- Note: Instead of listing all attributes, you can use '\*'  
to mean all attributes

-- Note 2: multiple relations mean cross product!

-- Note 3: distinct keyword?

# Comparison Operators

- Common comparison operators for the **WHERE** clause:



```
expr1 = expr2 (or expr1 == expr2 in SQLite)
```

```
expr1 < expr2
```

```
expr1 > expr2
```

```
expr1 <= expr2
```

```
expr1 >= expr2
```

```
expr1 <> expr2 (or expr1 != expr2 in SQLite)
```

```
attr IS [NOT] NULL
```

```
attr [NOT] BETWEEN expr1 AND expr2
```

```
attr [NOT] LIKE expr
```

```
attr [NOT] GLOB regexp (not supported by most other DBMS)
```

```
attr [NOT] IN expr (later)
```

```
attr [NOT] EXISTS (later)
```

# Let's Run Some SFW Queries

passengers

<u>f_name</u>	<u>m_name</u>	<u>l_name</u>	<u>ssn</u>
Homer	J	Simpson	111-11-1111
Bart	H	Simpson	444-44-4444
Lisa	G	Simpson	222-22-2222
Timothy		Lovejoy	555-55-5555
Joe	N	Quimby	666-66-6666
Ned	T	Flanders	777-77-7777
Frank		Ryerson	333-33-3333

onboard

<u>ssn</u>	<u>flight_no</u>	<u>seat</u>
555-55-5555	495	32 F
111-11-1111	86	1 D
777-77-7777	5932	25 A
666-66-6666	720	30 C
444-44-4444	5031	30 C
777-77-7777	303	25 A

plane

<u>tail_no</u>	<u>make</u>	<u>model</u>	<u>capacity</u>	<u>mph</u>
0	Boeing	747	525	570
1	Boeing	747	525	570
2	Airbus	A350	270	580
3	McDonnell Douglas	DC10	380	610

flight

<u>flight_no</u>	<u>dep_loc</u>	<u>dep_time</u>	<u>arr_loc</u>	<u>arr_time</u>	<u>tail_no</u>
720	Springfield, IL	7:15	Chicago, IL	7:45	3
86	Columbus, OH	16:00	Portland, OR	22:00	3
303	New York, NY	12:30	Miami, FL	13:00	1
1142	Paris, France	15:00	Munich, Germany	17:40	0
5932	Hartford, CT	9:00	Phoenix, AZ	12:00	0
495	Miami, FL	10:45	Austin, TX	13:30	1
5031	Akron, OH	14:20	Hartford, CT	16:45	2

```
sqlite> .read airport-schema.sql    -- creates the tables and constraints
sqlite> .read airport-populate.sql  -- populates the tables with above data
```

(Find these files on the course page!)

# Joins in SQL

## Nice-Looking Bowties



## Relational Algebra

$$R_1 \bowtie R_2$$

$$R_1 \bowtie_{\theta} R_2$$

$$R_1 \bowtie\!\!\!\bowtie R_2$$

$$R_1 \bowtie\!\!\!\lrcorner R_2$$

$$R_1 \bowtie\!\!\!\lrcorner\!\!\!\lrcorner R_2$$

## SQL Equivalent

$R_1$  NATURAL JOIN  $R_2$

*(next slide)*

$R_1$  LEFT OUTER NATURAL JOIN  $R_2$

$R_1$  RIGHT OUTER NATURAL JOIN  $R_2$

$R_1$  FULL OUTER NATURAL JOIN  $R_2$

*(Didn't I say these were "extended" relational operators?)*

# Inner-Join (Also known as Theta Join)

- ▶ The \theta-join is defined:  $\sigma_{\theta}(R_1 \times R_2) = R_1 \bowtie_{\theta} R_2$



- ▶ If you declare more than one relation in the [FROM](#) clause, it performs a cross product on all declared relations

$R_1 \times R_2$        $\leq$  same as  $\Rightarrow$       FROM [R\\_1](#), [R\\_2](#)

- ▶ The rest? Just use [WHERE](#) conditions to formulate the Join

SELECT \* FROM R1, R2 [WHERE](#) theta;

- ▶ Or, the inner-join syntax:


SELECT \* FROM R1 JOIN R2 [ON](#) theta;



# Rename Operator

## ► Syntax

- First form:  $\rho_{R_2}(R_1)$
- Meaning: Renames relation  $R_1$  to  $R_2$
- Second form:  $\rho_{(b_1, \dots, b_n)}(R_1)$
- Meaning: Renames  $R_1(a_1, \dots, a_n)$  to  $R_1(b_1, \dots, b_n)$
- Third form:  $\rho_{R_2}(b_1, \dots, b_n)(R_1)$
- Meaning: Renames  $R_1(a_1, \dots, a_n)$  to  $R_2(b_1, \dots, b_n)$



# Renaming Relational Tables (1st Form)

## ► Relational Algebra Syntax

- **First form:**  $\rho_{R_2}(R_1)$
- Meaning: Renames relation  $R_1$  to  $R_2$

## ► SQL syntax:

```
SELECT    ...  
FROM      R1 as X [ , R2 as Y, R3 as Z, ... ]  
[WHERE    C];  
  
-- Renames R_1 to A' [R_2 to B', ... ]  
-- R_1 can no longer be referred to in the query
```

```
sqlite> select * from passengers as A where A.l_name='Flanders';
```

f_name	m_name	l_name	ssn
Ned	T	Flanders	777-77-7777

```
sqlite> select * from passengers as A where passengers.l_name='Flanders';  
Error: no such column: passengers.l_name
```

# Renaming Attributes (2nd Form)

## ► Relational Algebra Syntax

- **Second form:**  $\rho_{(b_1, \dots, b_n)}(R_1)$
- Meaning: Renames  $R_1$ 's attributes from  $R_1(a_1, \dots, a_n)$  to  $R_1(b_1, \dots, b_n)$

## ► SQL syntax:

```
SELECT  a1 as b1 [a2 as b2, ...]  
FROM    ...  
[WHERE   C];  
  
-- Renames a1 to b1, a2 to b2, etc.
```

```
sqlite> select f_name as F, l_name as L from passengers where l_name='Simpson';
```

F	L
-----	-----
Homer	Simpson
Bart	Simpson
Lisa	Simpson

# Renaming Relations and Attributes (3rd Form)

## ► Relational Algebra Syntax

- **Third form:**  $\rho_{R_2(b_1, \dots, b_n)}(R_1)$
- **Meaning:** Renames  $R_1(a_1, \dots, a_n)$  to  $R_2(b_1, \dots, b_n)$

## ► SQL Syntax:

```
SELECT  a1 as b1 [a2 as b2, ...]
FROM    R1 as X [, R2 as Y, R3 as Z, ...]
[WHERE   C];
```

*-- But SQL has different take...*

```
sqlite> select mph as speed from plane as pl where pl.mph > 500;
```

```
speed
-----
570
570
580
610
```

Operationally, SQLite splits the 3rd-form into two parts...

- (1) Apply 1st-form immediately, then
- (2) Apply 2nd-form

$$\rho_{(speed)}(\Pi_{mph}(\sigma_{pl.mph > 500}(\rho_{pl}(planes))))$$

# Tricky Rename Example from Before

- ▶ Find all pairs of passengers that share last names:

- (Homer, Bart), (Lisa, Bart), (Lisa, Homer)

- ▶ Let's try...

```
select p1.f_name, p2.f_name
from passengers as p1, passengers as p2
where p1.l_name = p2.l_name;
```

f_name	f_name
Homer	Bart
Homer	Homer
Homer	Lisa
Bart	Bart
Bart	Homer
Bart	Lisa
Lisa	Bart
Lisa	Homer
Lisa	Lisa
Frank	Frank
Robert	Robert
Ned	Ned
Frank	Frank

*(Problem 1: Don't want the same first names)*

# Tricky Rename Example from Before (Cont.)

- Find all pairs of passengers that share last names:

- (Homer, Bart), (Lisa, Bart), (Lisa, Homer)

- Let's try...

```
select p1.f_name, p2.f_name
from   passengers as p1, passengers as p2
where  p1.l_name = p2.l_name
       and
       p1.f_name != p2.f_name;
```

f_name	f_name
Homer	Bart
Homer	Lisa
Bart	Homer
Bart	Lisa
Lisa	Bart
Lisa	Homer

*(Problem 2: Want the same combinations just once!)*

# Tricky Rename Example from Before (Cont.)

- Find all pairs of passengers that share last names:

- (Homer, Bart), (Lisa, Bart), (Lisa, Homer)

- Let's try...

```
select p1.f_name, p2.f_name
from   passengers as p1, passengers as p2
where  p1.l_name = p2.l_name
       and
       p1.f_name > p2.f_name;
```

f_name	f_name
Homer	Bart
Lisa	Bart
Lisa	Homer

*(When comparison operators like > and < are used with text, it means...?)*

# Imposing Order on Results

- ▶ Databases do not guarantee any ordering on the returned results
  - But we can impose an ordering using the following *optional* clause:



Any *SELECT ... FROM ... WHERE ...* query

SFW

```
[ORDER BY a_1 [DESC] [, a_2 [DESC], ..., a_k [DESC]]];
```

```
-- Sorts results by the given attribute(s) before returning
```

```
-- Sorts by a_1 first, then a_2, then a_3, ...
```

```
-- DESC keyword sorts results in descending order
```



# Imposing Order on Results (Cont.)

```
sqlite> select * from plane order by make, tail_no;
```

tail_no	make	model	capacity	mph
-----	-----	-----	-----	-----
3	Airbus	A350	270	580
5	Airbus	A380	200	500
1	Boeing	747	525	570
2	Boeing	747	525	570
4	McDonnell D	DC10	380	610

```
sqlite> select * from passengers order by l_name, f_name desc;
```

f_name	m_name	l_name	ssn
-----	-----	-----	-----
Ned	T	Flanders	777-77-7777
Frank	NULL	Lovejoy	555-55-5555
Robert	N	Quimby	666-66-6666
Frank	NULL	Ryerson	333-33-3333
Lisa	G	Simpson	222-22-2222
Homer	J	Simpson	111-11-1111
Bart	H	Simpson	444-44-4444

# Set Operators

- ▶ We've already seen cross-product through implicit joins, but what about these?
  - Remember, to apply any of these operations, the two relations must be *compatible*

- ▶ SQL Syntax:

- SFW1 UNION [ALL] SFW2
- SFW1 INTERSECT [ALL] SFW2
- SFW1 EXCEPT [ALL] SFW2

(ALL keyword retains duplicates!)

- ▶ Examples:

1. Get Lisa and Ned's passenger info
2. Get all last names and plane model numbers
3. Get everyone but Bart

# For Reference

## passengers

<u>f_name</u>	<u>m_name</u>	<u>l_name</u>	<u>ssn</u>
Homer	J	Simpson	111-11-1111
Bart	H	Simpson	444-44-4444
Lisa	G	Simpson	222-22-2222
Timothy		Lovejoy	555-55-5555
Joe	N	Quimby	666-66-6666
Ned	T	Flanders	777-77-7777
Frank		Ryerson	333-33-3333

## onboard

<u>ssn</u>	<u>flight_no</u>	<u>seat</u>
555-55-5555	495	32 F
111-11-1111	86	1 D
777-77-7777	5932	25 A
666-66-6666	720	30 C
444-44-4444	5031	30 C
777-77-7777	303	25 A

## plane

<u>tail_no</u>	<u>make</u>	<u>model</u>	<u>capacity</u>	<u>mph</u>
0	Boeing	747	525	570
1	Boeing	747	525	570
2	Airbus	A350	270	580
3	McDonnell Douglas	DC10	380	610

## flight

<u>flight_no</u>	<u>dep_loc</u>	<u>dep_time</u>	<u>arr_loc</u>	<u>arr_time</u>	<u>tail_no</u>
720	Springfield, IL	7:15	Chicago, IL	7:45	3
86	Columbus, OH	16:00	Portland, OR	22:00	3
303	New York, NY	12:30	Miami, FL	13:00	1
1142	Paris, France	15:00	Munich, Germany	17:40	0
5932	Hartford, CT	9:00	Phoenix, AZ	12:00	0
495	Miami, FL	10:45	Austin, TX	13:30	1
5031	Akron, OH	14:20	Hartford, CT	16:45	2

# Example with the [ALL] Keyword

```
sqlite> select l_name from passengers union  
        select model from plane;
```

```
l_name  
-----  
747  
A350  
A380  
DC10  
Flanders  
Lovejoy  
Quimby  
Ryerson  
Simpson
```

```
sqlite> select l_name from passengers union all  
        select model from plane;
```

```
l_name  
-----  
Simpson  
Simpson  
Simpson  
Lovejoy  
Quimby  
Flanders  
Ryerson  
747  
747  
A350  
DC10  
A380
```

# Topics

- ▶ Structured Query Language (SQL)
  - Data Definition Language (DDL)
  - Data Manipulation Language (DML)
    - Insert
    - Delete
    - Update
    - Select From Where
      - order by
      - joins (implicit, explicit)
      - set operations
        - » sub-queries
      - aggregation and grouping

# Subqueries

passengers

f_name	m_name	l_name	<u>ssn</u>
Homer	J	Simpson	111-11-1111
Bart	H	Simpson	444-44-4444
Lisa	G	Simpson	222-22-2222
Frank		Lovejoy	555-55-5555
Robert	N	Quimby	666-66-6666
Ned	T	Flanders	777-77-7777
Frank		Ryerson	333-33-3333

onboard

<u>ssn</u>	<u>flight no</u>	seat
555-55-5555	495	32 F
111-11-1111	86	1 D
777-77-7777	5932	25 A
666-66-6666	720	30 C
444-44-4444	5031	30 C
777-77-7777	303	25 A



► Recall that queries could be nested in relational algebra. We saw this example:

- Find the last names of all passengers not onboard any flights:

$$\Pi_{l\_name}(\sigma_{ssn \notin \Pi_{ssn}(onboard)}(passenger))$$

(1) Nested query returns the set of ssn for all passengers onboard any flight

Nested Projection Result =

<u>ssn</u>
555-55-5555
111-11-1111
777-77-7777
666-66-6666
444-44-4444
777-77-7777

# Subqueries (Cont.)

passengers

f_name	m_name	l_name	ssn
Homer	J	Simpson	111-11-1111
Bart	H	Simpson	444-44-4444
Lisa	G	Simpson	222-22-2222
Frank		Lovejoy	555-55-5555
Robert	N	Quimby	666-66-6666
Ned	T	Flanders	777-77-7777
Frank		Ryerson	333-33-3333

onboard

ssn	flight_no	seat
555-55-5555	495	32 F
111-11-1111	86	1 D
777-77-7777	5932	25 A
666-66-6666	720	30 C
444-44-4444	5031	30 C
777-77-7777	303	25 A



► Recall that queries could be nested in relational algebra. We saw this example:

- Find the last names of all passengers not onboard any flights:

$$\Pi_{l\_name}(\sigma_{ssn \notin \Pi_{ssn}(onboard)}(passenger))$$

(2) Now we select all the tuples from passengers where ssn is NOT IN the set returned by the nested query (1)

Select Result =

f_name	m_name	l_name	ssn
Lisa	G	Simpson	222-22-2222
Frank		Ryerson	333-33-3333

# Subqueries (Cont.)

passengers

f_name	m_name	l_name	ssn
Homer	J	Simpson	111-11-1111
Bart	H	Simpson	444-44-4444
Lisa	G	Simpson	222-22-2222
Frank		Lovejoy	555-55-5555
Robert	N	Quimby	666-66-6666
Ned	T	Flanders	777-77-7777
Frank		Ryerson	333-33-3333

onboard

ssn	flight_no	seat
555-55-5555	495	32 F
111-11-1111	86	1 D
777-77-7777	5932	25 A
666-66-6666	720	30 C
444-44-4444	5031	30 C
777-77-7777	303	25 A



► Recall that queries could be nested in relational algebra. We saw this example:

- Find the last names of all passengers not onboard any flights:


$$\Pi_{l\_name}(\sigma_{ssn \notin \Pi_{ssn}(onboard)}(passenger))$$

(3) Project lastname of all tuples selected in step (2).

Final Projection Result =

l_name
Simpson
Ryerson





# Subqueries: In, Not-in

- ▶ What was the key operator?
  - Checking to see if tuples were *in* or *not-in* the nested query result
- ▶ We can use these with nested queries in the WHERE clause

```
SELECT ...  
FROM ...  
WHERE a [NOT] IN (SWF);
```

*-- selects all tuples such that attribute a is in or not in the results of a nested SWF query*

```
sqlite> select l_name from passengers where ssn NOT IN (select ssn from onboard);
```

```
l_name  
-----  
Simpson  
Ryerson
```

# Subquery Example 2

passengers

<u>f_name</u>	<u>m_name</u>	<u>l_name</u>	<u>ssn</u>
Homer	J	Simpson	111-11-1111
Bart	H	Simpson	444-44-4444
Lisa	G	Simpson	222-22-2222
Frank		Lovejoy	555-55-5555
Robert	N	Quimby	666-66-6666
Ned	T	Flanders	777-77-7777
Frank		Ryerson	333-33-3333

onboard

<u>ssn</u>	<u>flight_no</u>	<u>seat</u>
555-55-5555	6	32F
111-11-1111	2	2B
222-22-2222	4	1F
777-77-7777	5	25A
333-33-3333	3	25A
666-66-6666	1	30C
444-44-4444	7	30C
777-77-7777	3	25A
555-55-5555	4	25A

plane

<u>tail_no</u>	<u>make</u>	<u>model</u>	<u>capacity</u>	<u>mph</u>
1	Boeing	747	525	570
2	Boeing	747	525	570
3	Airbus	A350	270	580
4	McDonnell Douglas	DC10	380	610
5	Airbus	A380	200	500

For SQL Subquery Lectures

Notes:

Airbus A380 not flying any flights

flight

<u>flight_no</u>	<u>dep_loc</u>	<u>dep_time</u>	<u>arr_loc</u>	<u>arr_time</u>	<u>tail_no</u>
1	Springfield, IL	7:15	Chicago, IL	7:45	4
2	Columbus, OH	16:00	Portland, OR	22:00	4
3	New York, NY	12:30	Miami, FL	13:00	2
4	Paris, France	15:00	Munich, Germany	17:40	1
5	Hartford, CT	9:00	Phoenix, AZ	12:00	1
6	Miami, FL	10:45	Austin, TX	13:30	2
7	Akron, OH	14:20	Hartford, CT	16:45	3

► Find the **f\_name**, **l\_name** of passengers sitting in 25A on any flight.

- Let's use both join and sub-query

# Subquery Example 3

passengers

<u>f_name</u>	<u>m_name</u>	<u>l_name</u>	<u>ssn</u>
Homer	J	Simpson	111-11-1111
Bart	H	Simpson	444-44-4444
Lisa	G	Simpson	222-22-2222
Frank		Lovejoy	555-55-5555
Robert	N	Quimby	666-66-6666
Ned	T	Flanders	777-77-7777
Frank		Ryerson	333-33-3333

onboard

<u>ssn</u>	<u>flight_no</u>	<u>seat</u>
555-55-5555	6	32F
111-11-1111	2	2B
222-22-2222	4	1F
777-77-7777	5	25A
333-33-3333	3	25A
666-66-6666	1	30C
444-44-4444	7	30C
777-77-7777	3	25A
555-55-5555	4	25A

plane

<u>tail_no</u>	<u>make</u>	<u>model</u>	<u>capacity</u>	<u>mph</u>
1	Boeing	747	525	570
2	Boeing	747	525	570
3	Airbus	A350	270	580
4	McDonnell Douglas	DC10	380	610
5	Airbus	A380	200	500

For SQL Subquery Lectures

Notes:

Airbus A380 not flying any flights

flight

<u>flight_no</u>	<u>dep_loc</u>	<u>dep_time</u>	<u>arr_loc</u>	<u>arr_time</u>	<u>tail_no</u>
1	Springfield, IL	7:15	Chicago, IL	7:45	4
2	Columbus, OH	16:00	Portland, OR	22:00	4
3	New York, NY	12:30	Miami, FL	13:00	2
4	Paris, France	15:00	Munich, Germany	17:40	1
5	Hartford, CT	9:00	Phoenix, AZ	12:00	1
6	Miami, FL	10:45	Austin, TX	13:30	2
7	Akron, OH	14:20	Hartford, CT	16:45	3

► Find the **f\_name** of passengers sitting in 25A on any flight.

- Can we still answer this query using only joins?

# Subquery Example 4

passengers

<u>f_name</u>	<u>m_name</u>	<u>l_name</u>	<u>ssn</u>
Homer	J	Simpson	111-11-1111
Bart	H	Simpson	444-44-4444
Lisa	G	Simpson	222-22-2222
Frank		Lovejoy	555-55-5555
Robert	N	Quimby	666-66-6666
Ned	T	Flanders	777-77-7777
Frank		Ryerson	333-33-3333

onboard

<u>ssn</u>	<u>flight_no</u>	<u>seat</u>
555-55-5555	6	32F
111-11-1111	2	2B
222-22-2222	4	1F
777-77-7777	5	25A
333-33-3333	3	25A
666-66-6666	1	30C
444-44-4444	7	30C
777-77-7777	3	25A
555-55-5555	4	25A

plane

<u>tail_no</u>	<u>make</u>	<u>model</u>	<u>capacity</u>	<u>mph</u>
1	Boeing	747	525	570
2	Boeing	747	525	570
3	Airbus	A350	270	580
4	McDonnell Douglas	DC10	380	610
5	Airbus	A380	200	500

For SQL Subquery Lectures

Notes:


Airbus A380 not flying any flights

flight

<u>flight_no</u>	<u>dep_loc</u>	<u>dep_time</u>	<u>arr_loc</u>	<u>arr_time</u>	<u>tail_no</u>
1	Springfield, IL	7:15	Chicago, IL	7:45	4
2	Columbus, OH	16:00	Portland, OR	22:00	4
3	New York, NY	12:30	Miami, FL	13:00	2
4	Paris, France	15:00	Munich, Germany	17:40	1
5	Hartford, CT	9:00	Phoenix, AZ	12:00	1
6	Miami, FL	10:45	Austin, TX	13:30	2
7	Akron, OH	14:20	Hartford, CT	16:45	3

► Find the planes that are not flying into Germany

- Join version, Subquery version



- In addition to in/not-in, we have exists/not-exist to test whether the subquery returns a set with elements or an empty set.

```
outer-SFW [NOT] EXISTS (inner-SFW);
```

*-- selects all tuples such that the nested SWF query returns anything (or nothing)*

*This is an expensive operation! Here's what it does:*

```
for each tuple t returned in the outer SFW,  
  run the inner SFW  
  if inner SFW [doesn't] returns anything //[NOT] EXISTS?  
    retain t  
  else  
    discard t
```

# Subquery Example 5

passengers

<u>f_name</u>	<u>m_name</u>	<u>l_name</u>	<u>ssn</u>
Homer	J	Simpson	111-11-1111
Bart	H	Simpson	444-44-4444
Lisa	G	Simpson	222-22-2222
Frank		Lovejoy	555-55-5555
Robert	N	Quimby	666-66-6666
Ned	T	Flanders	777-77-7777
Frank		Ryerson	333-33-3333

onboard

<u>ssn</u>	<u>flight_no</u>	<u>seat</u>
555-55-5555	6	32F
111-11-1111	2	2B
222-22-2222	4	1F
777-77-7777	5	25A
333-33-3333	3	25A
666-66-6666	1	30C
444-44-4444	7	30C
777-77-7777	3	25A
555-55-5555	4	25A

plane

<u>tail_no</u>	<u>make</u>	<u>model</u>	<u>capacity</u>	<u>mph</u>
1	Boeing	747	525	570
2	Boeing	747	525	570
3	Airbus	A350	270	580
4	McDonnell Douglas	DC10	380	610
5	Airbus	A380	200	500

For SQL Subquery Lectures

Notes:

Airbus A380 not flying any flights

flight

<u>flight_no</u>	<u>dep_loc</u>	<u>dep_time</u>	<u>arr_loc</u>	<u>arr_time</u>	<u>tail_no</u>
1	Springfield, IL	7:15	Chicago, IL	7:45	4
2	Columbus, OH	16:00	Portland, OR	22:00	4
3	New York, NY	12:30	Miami, FL	13:00	2
4	Paris, France	15:00	Munich, Germany	17:40	1
5	Hartford, CT	9:00	Phoenix, AZ	12:00	1
6	Miami, FL	10:45	Austin, TX	13:30	2
7	Akron, OH	14:20	Hartford, CT	16:45	3

► Find the fastest plane in the fleet

# Topics

- ▶ Structured Query Language (SQL)
  - Data Definition Language (DDL)
  - Data Manipulation Language (DML)
    - Insert
    - Delete
    - Update
    - Select
      - from, where
      - order by
      - set operations
      - joins (implicit, explicit)
      - sub-queries
      - aggregation and grouping

# Aggregation and Grouping

- Recall the R.A. Syntax  $g_1, \dots, g_j \mathcal{G} f_1(a_1), \dots, f_k(a_k)(R)$

Employees

<u>ENO</u>	Dept	Country	Name	Wage
0	A	US	John	50
1	A	China	Lynn	75
3	B	US	Ross	60
7	C	US	Julia	95
8	B	China	David	25
9	A	China	Ned	65

- Q1: Get the count of all employees
- Q2: Get the count of all employees, min wage, and max wage
- Q3: Get the count of all employees, min wage, and max wage by department and country

*(Let's do each in R.A. first)*



# Aggregation and Grouping in SQL

```

SELECT  f_1([distinct] a_1)[, ..., f_k([distinct] a_k),
                                g_1, ..., g_j,
                                a_1, ..., a_n]

FROM    R_1 [, R_2, ...]
WHERE   C
[GROUP BY g_1, g_2, ..., g_j [HAVING C']];

```

Aggregation  
Functions

Groups

Other  
attributes

Groups

Group cond.

```
sqlite> select avg(capacity), avg(mph) from plane group by make;
```

```
avg(capacity)  avg(mph)
```

```
-----
```

```
235.0          540.0
```

```
525.0          570.0
```

```
380.0          610.0
```

**\*\* What happened to my group, 'make'? \*\***  
**In SQL, you have to project the groups too**  
**(You did not have to do this with R.A.)**

```
sqlite> select avg(capacity), avg(mph), make from plane group by make;
```

```
avg(capacity)  avg(mph)  make
```

```
-----
```

```
235.0          540.0      Airbus
```

```
525.0          570.0      Boeing
```

```
380.0          610.0      McDonnell D
```

# HAVING Clause: Selecting Only Certain Groups

```

SELECT f_1([distinct] a_1)[, ..., f_k([distinct] a_k),
                                g_1, ..., g_j,
                                a_1, ..., a_n]
FROM R_1 [, R_2, ...]
WHERE C
[GROUP BY g_1, g_2, ..., g_j [HAVING C']];

```

Return only  
groups "Having"  
condition C'

- Find all departments, countries, and their average wages whose average wage is < \$50
- Find the department that pays the highest average wage

Employees

<u>ENO</u>	Dept	Country	Name	Wage
0	A	US	John	50
1	A	China	Lynn	75
3	B	US	Ross	60
7	C	US	Julia	95
8	B	China	David	25
9	A	China	Ned	65

# Topics

- ▶ Structured Query Language (SQL)
  - Data Definition Language (DDL)
    - Create table
    - Drop table
    - Alter table
  - Data Manipulation Language (DML)
    - SFW, Rename, Joins
    - Ordering results
    - Subqueries
    - Grouping/Aggregation
  - Views
- ▶ Conclusion

# Views

- ▶ Often, it is undesirable for all users to see entire database
- ▶ For example, security and privacy concerns...
  - Health records (HIPAA)
  - Student records (FERPA)
- ▶ More desirable:
  - We only want some users to see parts of the database that certain users are authorized to see.



# Views (Cont.)

- ▶ Give certain users only a **read-only** *logical view* of the entire database!
  - Users cannot use write operations on the view!
  - insert, delete, update prohibited!
- ▶ Syntax:

```
create view V as SFW;
```

```
drop view V;
```

# Example of Views in SQLite

- ▶ Hide passenger SSN from flight attendants
- ▶ Show only first/last name, flight number, and seat assignment
- ▶ Sorted by last name
- ▶ Call this view **passinfo**

passengers

f_name	m_name	l_name	ssn
Homer	J	Simpson	111-11-1111
Bart	H	Simpson	444-44-4444
Lisa	G	Simpson	222-22-2222
Frank		Lovejoy	555-55-5555
Robert	N	Quimby	666-66-6666
Ned	T	Flanders	777-77-7777
Frank		Ryerson	333-33-3333

onboard

ssn	flight_no	seat
555-55-5555	495	32 F
111-11-1111	86	1 D
777-77-7777	5932	25 A
666-66-6666	720	30 C
444-44-4444	5031	30 C
777-77-7777	303	25 A

passinfo

f_name	l_name	flight_no	seat
Ned	Flanders	5932	25 A
Ned	Flanders	303	25 A
Frank	Lovejoy	495	32 F
Robert	Quimby	720	30 C
Bart	Simpson	5031	30 C
Homer	Simpson	86	1 D

The View

# Syntax: Creating Views

passengers

<u>f_name</u>	<u>m_name</u>	<u>l_name</u>	<u>ssn</u>
Homer	J	Simpson	111-11-1111
Bart	H	Simpson	444-44-4444
Lisa	G	Simpson	222-22-2222
Frank		Lovejoy	555-55-5555
Robert	N	Quimby	666-66-6666
Ned	T	Flanders	777-77-7777
Frank		Ryerson	333-33-3333

onboard

<u>ssn</u>	<u>flight_no</u>	<u>seat</u>
555-55-5555	495	32 F
111-11-1111	86	1 D
777-77-7777	5932	25 A
666-66-6666	720	30 C
444-44-4444	5031	30 C
777-77-7777	303	25 A

passinfo

<u>f_name</u>	<u>l_name</u>	<u>flight_no</u>	<u>seat</u>
Ned	Flanders	5932	25 A
Ned	Flanders	303	25 A
Frank	Lovejoy	495	32 F
Robert	Quimby	720	30 C
Bart	Simpson	5031	30 C
Homer	Simpson	86	1 D

## ► SQL:

```
create view passinfo as
  select f_name, l_name, flight_no, seat from passengers
  natural join onboard order by l_name;
```

# Views in SQLite

```
sqlite> create view passinfo as select f_name,l_name,flight_no,seat
from passengers natural join onboard order by l_name;
```

```
sqlite> select * from passinfo;
```

f_name	l_name	flight_no	seat
-----	-----	-----	-----
Ned	Flanders	3	25A
Ned	Flanders	5	25A
Frank	Lovejoy	6	32F
Robert	Quimby	1	30C
Homer	Simpson	2	2B
Bart	Simpson	7	30C

```
sqlite> delete from passinfo where l_name='Simpson';
```

```
Error: cannot modify passinfo because it is a view
```



# Topics

- ▶ Structured Query Language (SQL)
  - Data Definition Language (DDL)
    - Create table
    - Drop table
    - Alter table
  - Data Manipulation Language (DML)
    - SFW, Rename, Joins
    - Ordering results
    - Subqueries
    - Grouping/Aggregation
- ▶ Conclusion

# In Conclusion...

- ▶ SQL first appeared in **IBM System-R** (1976)
  - Now the standard relational data query language
  - Same expressivity as relational algebra, but user-friendly
- ▶ Two parts:
  - DDL: Deals with structure of database
  - DML: Deals with data
- ▶ Next: What constitutes **good DB design**?
  - Entity-Relationship (ER) Model
  - Normalization

